



## **GPGPU**

**General Purpose computation on Graphics Processing Units**

**Mark Harris, 2002**

**Perform demanding calculations on the GPU instead of the CPU!**

**At first, appeared to be a wild idea, but is now a very serious technology! Results were highly varied in the early years, but the GPU advantage has grown bigger and bigger.**



## **Key components of the GPGPU trend**

High processing power in parallel

Programmability: Introduction of shader programs, much more flexible, programmable for any problem.

Floating-point buffers: Vital! Initially with poor precision. 32-bit floating-point decent... but not really impressive.



## **GPGPU solutions**

- Using fixed pipeline graphics
  - Shader programs
    - CUDA
    - OpenCL



## **Fixed pipeline GPGPU**

Reformulate a problem to something that can be done by standard graphics operations.

Limited success 1999/2000. Not of any practical interest!



## **Shader-based GPGPU**

**Portable! Most GPUs can use shaders, no need for extra software, run using standard software/drivers.**

**All modern shader languages (GLSL, Cg, HLSL) are similar and easy to program in.**

**Requires a re-mapping of data to textures.**

**Very good results already in 2005: 8x speedups overall reported!**



## **GPGPU using shaders**

**Has less attention now, due to CUDA.**

**Still interesting:**

- **Apart from some reusable standard code, it is not very complicated.**
- **Portable to most GPUs with no extra software.**
- **Excellent performance.**



Information Coding / Computer Graphics, ISY, LiTH

## **CUDA-based GPGPU**

Only works on NVidia hardware.

Requires extra software - which isn't very elegant.

Nice integration of CPU and GPU code in the same program.

Excellent results! 100x speedups are common - before optimizing! Even low-end GPUs give significant boosts.



Information Coding / Computer Graphics, ISY, LiTH

## **OpenCL-based GPGPU**

Works on various hardware - not only GPUs.

Developed by Khronos Group, pushed by Apple.

Harder to get started, software looks pretty much like programming shaders.



## Use the source, Luke!

Three trivial examples:

Hello World! for CUDA

Hello World! for OpenCL

Hello World for GLSL



## So what GPU should you get?

For CUDA, go for Fermi or Kepler boards!

Kepler has more cores but Fermi is still strong in double precision.

GTX560 good middle-range board, best Fermi price/performance and reasonable power consumption.

GTX660Ti best Kepler price/performance and even better power consumption.

Avoid overclocked boards!

Don't bother with "professional" Quadro boards.

AMD is DEFINITELY an option for shaders or OpenCL but not CUDA.



Information Coding / Computer Graphics, ISY, LiTH

## **In the Southfork lab**

**GTX660Ti**

**Exactly one year old. Still one of the fastest.**

**1300 cores!**

**Close-to-high-end mid-range board. Great price/performance,  
lots of parallelism to play with, and pretty nice power  
consumption.**



Information Coding / Computer Graphics, ISY, LiTH

**That's all, folks!**

**Next friday: Introduction to CUDA**