



Information Coding / Computer Graphics, ISY, LiTH

## **Lecture 11 (#3 on GPU Computing)**

# **More CUDA**



## **In this episode...**

- **Query device capabilities**
  - **CUDA events**
- **More on CUDA memory:**


**Coalescing, Constant memory, Texture memory...**



## Lab 4

**First version of the new lab available since yesterday (sunday)**

**Major change: "Mandelbrot revisited" part, to follow up lab 1.**





## **The story so far...**

- **CUDA and its language extensions**
  - **The CUDA architecture**
    - **Intro to memory**
- **Matrix multiplication example, using shared memory**



## **CUDA and its language extensions**

**Kernel involation myKernel<<<>>>()**

**\_\_global\_\_ \_\_device\_\_ \_\_host\_\_**

**cudaMalloc(), cudaMemcpy()**

**threadIdx, blockIdx, blockDim, gridDim**

**Using nvcc**



**Information Coding / Computer Graphics, ISY, LiTH**

# **The CUDA architecture**

**Blocks and threads**

**Grid-block-thread hierarchy**

**Indexing data with thread/block numbers**



# **Intro to memory**

**global memory**

**shared memory**


**constant memory**

**local memory**

**texture memory/texture units**



## Matrix multiplication example, using shared memory



**Huge speedup - my measly 9400M went from obvious loser to clearly faster than CPU!**






**Information Coding / Computer Graphics, ISY, LiTH**

## **CUDA and graphics**

**Simplest way: Pass output from CUDA, typically to an OpenGL texture.**

**Example: Julia set.**


**Good for visualizing results. Better methods exist, without having to move data to CPU and back.**





## Self-squaring fractals, the Julia set

$$z_{k+1} = z_k^2 + \lambda$$



**Julia set for**  
 $\lambda = (0, 1) = 0 + j$

**Start with position in  
complex space.**

**Apply complex function  
recursively**

**Inspect distance to origin**

**Perfectly parallel  
algorithm**



**Information Coding / Computer Graphics, ISY, LiTH**

**Over to today's episode:**



## **Lecture questions:**

- 1. Why can using constant memory improve performance?**
- 2. What is CUDA Events used for?**
- 3. What does coalescing mean and what should we do to get a speedup from coalescing?**



## Query devices

**You can't trust all devices to have the same - or even similar - data.**

**New boards may have totally different data.**

**Query CUDA for a list of features using  
`cudaGetDeviceProperties()`**



## Example query result

```
----- Information for GeForce 9400M -----  
      Compute capability:  1.1  
Total global memory (VRAM): 259712 kB  
      Total constant Mem:  64 kB  
      Number of SMs:      2  
      Shared mem per SM:  16 kB  
      Registers per SM:   8192  
      Threads in warp:    32  
      Max threads per block: 512  
Max thread dimensions: (512, 512, 64)  
Max grid dimensions: (65535, 65535, 1)
```



# What is important?

Compute capability - can this board at all work with our program?

Amount of shared memory - make sure we fit.

Max threads, max dimensions - make sure we fit.

Threads in warp: A lower bound for performance.

Number of SMs: Lower bound for blocks



# Compute capability

**Essentially CUDA/architecture version number.**

- 1.0: Original release.**
- 1.1: Mapped memory, atomic operations.**
- 1.3: Double support.**
- 2.0: Fermi.**
- 3.0: Kepler.**
- 5.0: Maxwell.**





Feature Support	Compute Capability					
	1.0	1.1	1.2	1.3	2.x, 3.0	3.5
(Unlisted features are supported for all compute capabilities)						
Atomic functions operating on 32-bit integer values in global memory (Atomic Functions)	No	Yes				
atomicExch() operating on 32-bit floating point values in global memory (atomicExch())						
Atomic functions operating on 32-bit integer values in shared memory (Atomic Functions)	No	Yes				
atomicExch() operating on 32-bit floating point values in shared memory (atomicExch())						
Atomic functions operating on 64-bit integer values in global memory (Atomic Functions)						
Warp vote functions (Warp Vote Functions)						
Double-precision floating-point numbers	No			Yes		
Atomic functions operating on 64-bit integer values in shared memory (Atomic Functions)	No				Yes	
Atomic addition operating on 32-bit floating point values in global and shared memory (atomicAdd())						
__ballot() (Warp Vote Functions)						
__threadfence_system() (Memory Fence Functions)						
__syncthreads_count(), __syncthreads_and(), __syncthreads_or() (Synchronization Functions)						
Surface functions (Surface Functions)						
3D grid of thread blocks						
Funnel shift (see reference manual)	No					Yes

ITH



## Information Coding / Computer Graphics, ISY, LiTH

	FERMI GF100	FERMI GF104	KEPLER GK104	KEPLER GK110
<b>Compute Capability</b>	2.0	2.1	3.0	3.5
<b>Threads / Warp</b>	32	32	32	32
<b>Max Warps / Multiprocessor</b>	48	48	64	64
<b>Max Threads / Multiprocessor</b>	1536	1536	2048	2048
<b>Max Thread Blocks / Multiprocessor</b>	8	8	16	16
<b>32-bit Registers / Multiprocessor</b>	32768	32768	65536	65536
<b>Max Registers / Thread</b>	63	63	63	255
<b>Max Threads / Thread Block</b>	1024	1024	1024	1024
<b>Shared Memory Size Configurations (bytes)</b>	16K	16K	16K	16K
	48K	48K	32K	32K
			48K	48K
<b>Max X Grid Dimension</b>	$2^{16}-1$	$2^{16}-1$	$2^{32}-1$	$2^{32}-1$
<b>Hyper-Q</b>	No	No	No	Yes
<b>Dynamic Parallelism</b>	No	No	No	Yes

Compute Capability of Fermi and Kepler GPUs



## Information Coding / Computer Graphics, ISY, LiTH

<b>Compute Capability</b>	<b>1.0</b>	<b>1.1</b>	<b>1.2</b>	<b>1.3</b>	<b>2.0</b>	<b>2.1</b>	<b>3.0</b>	<b>3.5</b>
<i>SM Version</i>	sm_10	sm_11	sm_12	sm_13	sm_20	sm_21	sm_30	sm_35
<i>Threads / Warp</i>	32	32	32	32	32	32	32	32
<i>Warps / Multiprocessor</i>	24	24	32	32	48	48	64	64
<i>Threads / Multiprocessor</i>	768	768	1024	1024	1536	1536	2048	2048
<i>Thread Blocks / Multiprocessor</i>	8	8	8	8	8	8	16	16
<i>Max Shared Memory / Multiprocessor (bytes)</i>	16384	16384	16384	16384	49152	49152	49152	49152
<i>Register File Size</i>	8192	8192	16384	16384	32768	32768	65536	65536
<i>Register Allocation Unit Size</i>	256	256	512	512	64	64	256	256
<i>Allocation Granularity</i>	block	block	block	block	warp	warp	warp	warp
<i>Max Registers / Thread</i>	124	124	124	124	63	63	63	255
<i>Shared Memory Allocation Unit Size</i>	512	512	512	512	128	128	256	256
<i>Warp allocation granularity</i>	2	2	2	2	2	2	4	4
<i>Max Thread Block Size</i>	512	512	512	512	1024	1024	1024	1024
<i>Shared Memory Size Configurations (bytes)</i>	16384	16384	16384	16384	49152	49152	49152	49152
<i>[note: default at top of list]</i>					16384	16384	16384	16384
							32768	32768
<i>Warp register allocation granularities</i>					64	64	256	256
<i>[note: default at top of list]</i>					128	128		



# **Do I care about Compute capability?**

**While learning CUDA - not much. Stick to the basics, it works on all.**

**But if you write professional CUDA code, of course.**



# CUDA Events

## Timing!

**Two ways of timing CUDA programs:**

- **CPU timer. Synchronize at start and end.**
- **CUDA Events. Synchronize at end.**

**Synchronize? Because CUDA runs asynchronously.**



# CUDA Events API

**cudaEventCreate** - initialize an event variable

**cudaEventRecord** - place a marker in the queue

**cudaEventSynchronize** - wait until all markers  
have received values

**cudaEventElapsedTime** - get the time difference  
between two events



# **CUDA memory**

**Coalescing**

**Constant memory**

**Texture memory**

**Pinned memory**



# CUDA memory

We already know...

- Global memory is slow.
- Shared memory is fast and can be used as "manual cache"
- There were some other kinds of memory...







# Coalescing

**Always access global memory "in order"**

**If threads access global memory in order of thread numbers, performance will be improved!**



**Good!**



**Bad!**



# WTF?

**How can performance depend on what order  
I access my data??? Isn't it "random  
access"?**


**Yes... You can access in any order you want,  
but ordered access *helps* the GPU to read  
more data in one access!**

**Why? Because the GPU bus is wider than  
your data!**




# Coalescing

**Example: Assume that the data below is 1/4 of the bus width.**



**Good!**



**Bad!**



# **Coalescing on Fermi & later**

**Effect reduced by caches - but not removed.**

**Coalescing is still needed for maximum performance.**



## **Accelerating by coalescing**

**Pure memory transfers can be 10x faster by taking advantage of memory coalescing!**

**Example: Matrix transpose**

**No computations!**

**Only memory accesses.**



# Matrix transpose

## Naive implementation

```
__global__ void transpose_naive(float *odata, float* idata, int width, int height)
{
    unsigned int xIndex = blockDim.x * blockIdx.x + threadIdx.x;
    unsigned int yIndex = blockDim.y * blockIdx.y + threadIdx.y;


    if (xIndex < width && yIndex < height)
    {
        unsigned int index_in  = xIndex + width * yIndex;
        unsigned int index_out = yIndex + height * xIndex;
        odata[index_out] = idata[index_in];
    }
}
```

**How can this be bad?**



## Matrix transpose

Coalescing problems




Row-by-row and column-by-column.  
Column accesses non-coalesced!



## Matrix transpose

### Coalescing solution



**Read from global memory  
to shared memory**

**In order from global, any  
order to shared**

**Write to global memory**

**In order write to global,  
any order from shared**





# Better CUDA matrix transpose kernel

```
__global__ void transpose(float *odata, float *idata, int width, int height)
{
    __shared__ float block[BLOCK_DIM][BLOCK_DIM+1];

    // read the matrix tile into shared memory
    unsigned int xIndex = blockIdx.x * BLOCK_DIM + threadIdx.x;
    unsigned int yIndex = blockIdx.y * BLOCK_DIM + threadIdx.y;
    if((xIndex < width) && (yIndex < height))
    {
        unsigned int index_in = yIndex * width + xIndex;
        block[threadIdx.y][threadIdx.x] = idata[index_in];
    }

    __syncthreads();

    // write the transposed matrix tile to global memory
    xIndex = blockIdx.y * BLOCK_DIM + threadIdx.x;
    yIndex = blockIdx.x * BLOCK_DIM + threadIdx.y;
    if((xIndex < height) && (yIndex < width))
    {
        unsigned int index_out = yIndex * height + xIndex;
        odata[index_out] = block[threadIdx.x][threadIdx.y];
    }
}
```

Shared memory  
for temporary  
storage

Read data to  
temporary buffer

Write data to  
tglobal memory



## **Coalescing rules of thumb**

- **The data block should start on a multiple of 64**
- **It should be accessed in order (by thread number)**
- **It is allowed to have threads skipping their item**
  - **Data should be in blocks of 4, 8 or 16 bytes**




# Shared memory

**Split into multiple memory banks (32). Fastest if you access different banks with each thread**

**Interleaved, 32 bits chunks**

**Thus: Address in 32-bit steps between threads for best performance**





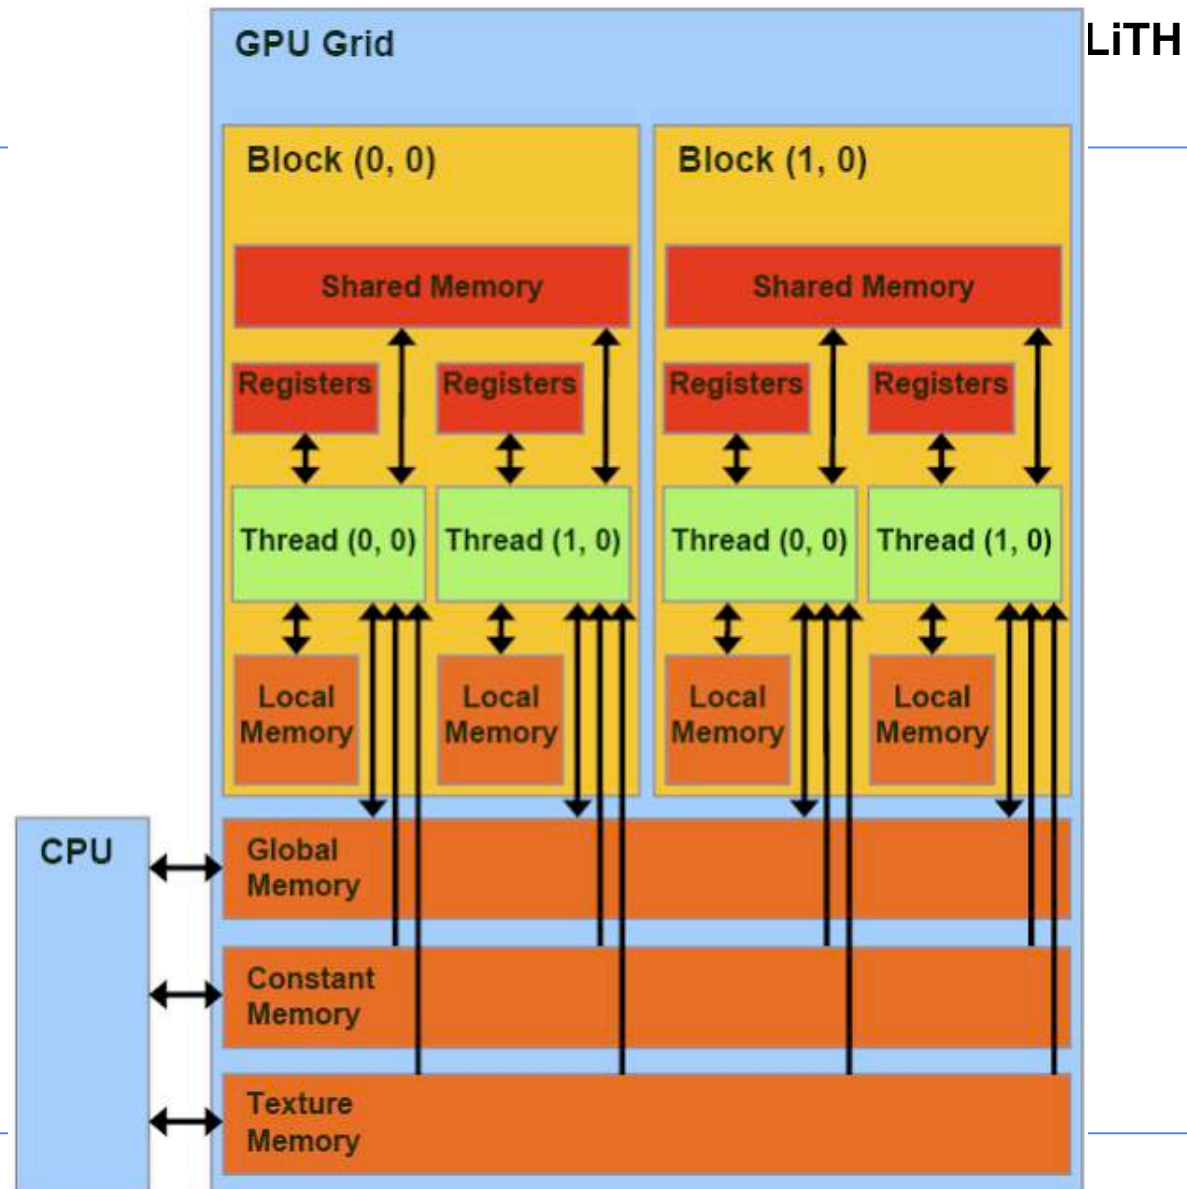
# Constant memory

Sounds boring... but has its uses.

Read-only (for kernels)

`__constant__` modifier

Use for input data, obviously





## Benefits of constant memory

- **No cudaMemcpy needed! Just use it from kernel, write from CPU!**
- **For data read by all threads, significantly faster than global memory!**
- **Read-only memory is easy to cache.**



## **Why faster access? When?**

**All threads reading the same data.**

**One read can be broadcast to all "nearby" threads.**

**Nearby? All threads in same "half-warp" (16 threads in most pre-Fermi architectures)**


**But no help if threads are reading different data!**



## Example of using constant memory: Ray-caster

Demo from "CUDA by example"

With and without using `__const__`







## Ray-caster example

Every thread renders one pixel

Loop through all spheres, find closest with intersection

Write result to an image buffer.

Image buffer displayed with OpenGL.

Non-const: Uploads sphere array by cudaMemcpy()

Const: Declares array `__const__`, uses directly from kernel.  
(Slightly simpler code!)



## **Ray-caster example**

**Resulting time:**

**Without using const: 70.2 ms**

**With const: 41.9 ms**

**Significant difference - for something that  
simplified the code!**



## **Constant memory conclusions**

**Relatively fast memory - for the case when all threads read the same memory!**

**Some advantage for code complexity.**

**NOT something we use for everything.**



## **Texture memory/ Texture units**

**Texture memory, yet another kind of memory (or  
memory access method)**

**But didn't we hide the graphics heritage...?**

**Access global memory though the texturing units.  
Lets CUDA take advantage of the strong points  
with texturing units.**



# Texture memory

**Read-only.**

**Cached! Can be fast if data access patterns are good.**

**Texture filtering, linear interpolation.**

**Especially good for handling 4 floats at a time (float4).**

**`cudaBindTextureToArray()` binds data to a texture unit.**



# Texture memory for graphics

Texture data mostly for rendering textures


One texel used by 4 neighbor pixels

One pixel usually 4 bytes - more than one pixel can be read on one read.

Designed for *spatial locality*



## Varying access patterns - but neighbors are still neighbors!





Information Coding / Computer Graphics, ISY, LiTH

# **Spatial locality for other things than textures**

**Image filters of local nature**

**Physics simulations with local updates, transfer of heat, liquids, pressure...**

**Big jumps, no gain!**





# Using texture memory in CUDA

**Allocate with cudaMalloc**

**Bind to texture unit using cudaBindTexture2D()**

**Read from data using tex2D()**

**Drawback: Just like in OpenGL, messy to keep track of which texture unit/texture reference is which data.**



# Interpolation

## Computation tricks when optimizing


Texture access provides hardware accelerated linear interpolation!

Access texture data on non-integer coordinates and the texture hardware will do linear interpolation automatically!

Can be used for many calculations, e.g. filters.



# Interpolation



**Texture accesses and calculations hardware accelerated!**



## **Hardware interpolation too good to be true...**

**The interpolation trick sounds kind of useful (for some cases)... but isn't as useful as it seems.**

**Why? It is ment for interpolating between texels, visually. Small errors is not a problem then! May have low precision, like 10 steps.**


**Not as fun then...**





# Demo using texture memory

## Heat transfer demo






# Demo using texture memory

## Heat transfer demo

**Makes local operations modelling heat dissipation**



**Seriously... pretty slow. I could beat this with pure OpenGL any time. Why?**



**Information Coding / Computer Graphics, ISY, LiTH**

**That's all folks!**

**Next: Sorting on the GPU.**