



Information Coding / Computer Graphics, ISY, LiTH

Lecture 13

OpenCL

GPU computing with GLSL

OpenGL Compute shaders



Lecture questions

- 1) What kind of devices will OpenCL run on?
- 2) What does an OpenCL work group correspond to in CUDA?
- 3) What geometry is typically used for shader-based GPU computing?
- 4) What are the main advantages of using OpenGL Compute Shaders?



Information Coding / Computer Graphics, ISY, LiTH

Introduction to OpenCL

Open Compute Language





- **Motivation**
- **Overview**
- **Examples**
- **Performance comparison**



Information Coding / Computer Graphics, ISY, LiTH

Origins of OpenCL

Initiated by Apple

Managed by Khronos group

Many supporting parties

Many providers





Why?

- The market could not let CUDA rule the world
 - Support for other platforms
 - Open standard
 - Similarity with OpenGL

For programming "all" parallel architectures



Information Coding / Computer Graphics, ISY, LiTH

Supported architectures

GPU

Intel compatible CPUs

ARM

FPGA

CELL

**Who decides? Any company making its own
OpenCL implementation!**



Information Coding / Computer Graphics, ISY, LiTH

”Open”?

Means *open specification*

Like OpenGL

Many providers making their own implementation

There is not *one* OpenCL library.



No free lunch

Model does not fit all architectures

**One size fits all - platform dependent
optimizations hard to do**



Information Coding / Computer Graphics, ISY, LiTH

OpenCL for GPU Computing

Mostly similar to CUDA both in architecture and performance!

Messy setup- but you get used to it

Kernels similar to CUDA

Easier for NVidia to be first with new features



OpenCL vs CUDA terminology

OpenCL

compute unit
work item
work group
local memory
private memory

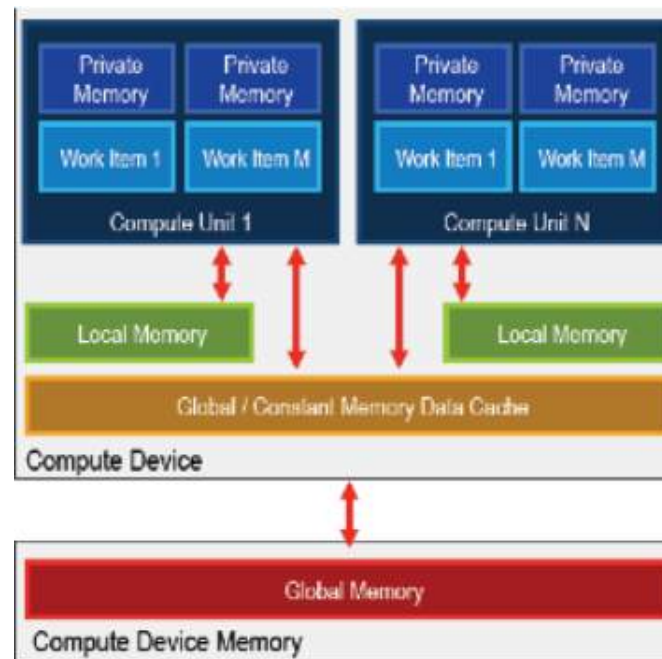
CUDA

multiprocessor (SM)
thread
block
shared memory
registers

And CUDA local memory =?
OpenCL local memory (= CUDA shared memory)



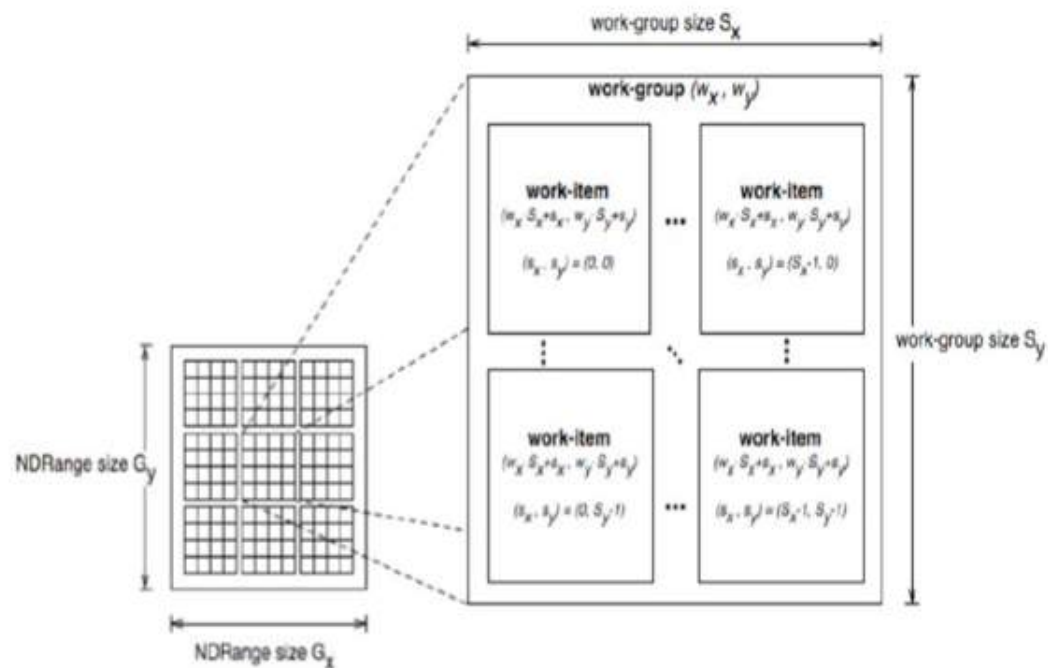
OpenCL memory model



Been there, done that...



OpenCL execution model



Anyone who see "blocks" and "threads"?



Synchronization

Kernels can synchronize within a work group:

```
barrier( CLK_LOCAL_MEM_FENCE )
```

No synchronization between work groups.

Synchronizes memory access. You choose which kind of memory access to synchronize (global, local).



Synchronization

The host (CPU) can synchronize on global level:

Available for:
tasks (e.g. `clEnqueueNDRangeKernel`)
Memory (e.g. `clEnqueueReadBuffer`)
events (e.g. `clWaitforEvents`)



Heterogenous

Some differences from CUDA: Designed for heterogenous systems!

Several devices may be active at once

You can specify which device to launch a task to

Query devices and device characteristics

Some overhead compared to CUDA, and the reward is flexibility!



Language

Based on C99, but:

- ❑ No function pointers
- ❑ No pointers to pointers in function calls
(=> no multi-dimensional arrays)
- ❑ No recursion
- ❑ No arrays with dynamical length
- ❑ No bitfields
- ❑ Also, no possibility to call a kernel from another kernel

Optional:

- ❑ Pointers with length <32 bit
- ❑ Writing support for 3D images
- ❑ Double and half types
- ❑ Atomic functions



Information Coding / Computer Graphics, ISY, LiTH

On the positive side:

- ❑ Integrated functions for reading / writing 2D images and reading 3D images
- ❑ Converting functions incl. explicit rounding and saturation
- ❑ math.h, all functions with different precisions
- ❑ Vector support (2-, 3- and 4-dimensional)

Available primitive datatypes:

- ❑ Bool, char, int, long, float, size_t, void, +unsigned versions

Mix of OpenCL and OpenGL possible

- ❑ Can share data structures and variables (without copying)
- ❑ API functions available



How about that setup?

- 1) Get a list of platforms**
- 2) Choose a platform**
- 3) Get a list of devices**
- 4) Choose a device**
- 5) Create a context**
- 6) Load and compile kernel code**



Then we can start working

- 7) Allocate memory**
- 8) Copy data to device**
- 9) Run kernel**
- 10) Wait for kernel to complete**
- 11) Read data from device**
- 12) Free resources**



1-5: Where to run

```
cl_platform_id platform;  
unsigned int no_plat;  
err = clGetPlatformIDs(1, &platform, &no_plat);
```

**Simplified here -
might fail!**

```
// Where to run  
err = clGetDeviceIDs(platform, CL_DEVICE_TYPE_GPU, 1, &device_id, NULL);  
if (err != CL_SUCCESS) return -1;
```

Context

```
context = clCreateContext(0, 1, &device_id, NULL, NULL, &err);  
if (!context) return -1;  
commands = clCreateCommandQueue(context, device_id, 0, &err);  
if (!commands) return -1;
```



6: Kernel

```
// What to run
program =
clCreateProgramWithSource(context, 1, (const
char **) & KernelSource, NULL, &err);
if (!program) return -1;

err = clBuildProgram(program, 0, NULL, NULL,
NULL, NULL);
if (err != CL_SUCCESS) return -1;
kernel = clCreateKernel(program, "hello", &err);
if (!kernel || err != CL_SUCCESS) return -1;
```

```
const char *KernelSource = "\n" \
"__kernel void hello(      \n" \
"    __global char* a,      \n" \
"    __global char* b,      \n" \
"    __global char* c,      \n" \
"    const unsigned int count) \n" \
"{                          \n" \
"    int i = get_global_id(0); \n" \
"    if(i < count)           \n" \
"        c[i] = a[i] + b[i]; \n" \
"}                          \n" \
"\n";
```

Most programs also load kernels from files



7-8: Get the data in there

```
// Create space for data and copy a and b to device (note that we could
also use clEnqueueWriteBuffer to upload)
input = clCreateBuffer(context, CL_MEM_READ_ONLY |
CL_MEM_USE_HOST_PTR, sizeof(char) * DATA_SIZE, a, NULL);
input2 = clCreateBuffer(context, CL_MEM_READ_ONLY |
CL_MEM_USE_HOST_PTR, sizeof(char) * DATA_SIZE, b, NULL);
output = clCreateBuffer(context, CL_MEM_WRITE_ONLY, sizeof(char) *
DATA_SIZE, NULL, NULL);
if (!input || !output) return -1;

// Send data
err = clSetKernelArg(kernel, 0, sizeof(cl_mem), &input);
err |= clSetKernelArg(kernel, 1, sizeof(cl_mem), &input2);
err |= clSetKernelArg(kernel, 2, sizeof(cl_mem), &output);
err |= clSetKernelArg(kernel, 3, sizeof(unsigned int), &count);
if (err != CL_SUCCESS) return -1;
```




9-10: Run kernel, wait for completion

```
// Run kernel!  
err = clEnqueueNDRangeKernel(commands, kernel, 1,  
    NULL, &global, &local, 0, NULL, NULL);  
  
if (err != CL_SUCCESS) return -1;  
  
clFinish(commands);
```



11-12: Read back data, release

```
// Read result
err = clEnqueueReadBuffer( commands, output, CL_TRUE, 0,
sizeof(char) * count, c, 0, NULL, NULL );
if (err != CL_SUCCESS) return -1;

// Print result
printf("%s\n", c);

// Clean up
clReleaseMemObject(input);
clReleaseMemObject(output);
clReleaseProgram(program);
clReleaseKernel(kernel);
clReleaseCommandQueue(commands);
clReleaseContext(context);
```



Information Coding / Computer Graphics, ISY, LiTH

"Platform" vs "device"

Platform = an OpenCL implementation

Device = a chip which the platform supports



Language freedom... sort of

- + Very easy to call from any language! Anything that can call into a C API can call OpenCL!**
- Kernel code is only C-style (although a specific implementation may choose to support more).**



Performance

Investigations report remarkably small differences

Our research on FFT so far has CUDA up to 2x faster

Very hard to compare, due to multiple OpenCL implementations

Some report CUDA to be better on NVidia platforms... some report a draw even there.

Our experience: Usually very close!



Conclusions on OpenCL

Don't fear the complex setup phase! The rest is similar to CUDA.

Performance tend to be on par with CUDA or almost.

Speciality: heterogenous systems!