



## Parallel Sorting Contest 2014

The winners are ...

Christoph Kessler, IDA,  
Linköpings universitet

## Parallel Sorting Contest 2014



### ■ Category Konrad Zuse:

- 
- 

### ■ Category Southfork:

- 
- 



... and thanks to all who participated!

C. Kessler, IDA, Linköpings universitet.

2



## Wrap-Up and Outlook

Christoph Kessler, IDA,  
Linköpings universitet

\* Similar as in TDDC78



## Lectures (1)

- **Lecture 1:** Motivation, Multicore architectural concepts and trends.
- **Lecture 2:** Parallel programming with threads and tasks.
- **Lesson 1:** How to measure and visualize performance of parallel programs. CPU lab introduction.
- **Lecture 3:** Shared memory architecture concepts and performance issues\*.
- **Lecture 4:** Non-blocking synchronization.
- **Lecture 5-6:** Theory: Design and analysis of parallel algorithms\*.
- **Lecture 7:** Parallel sorting algorithms.
- **Lesson 2:** Selected theory exercises.

C. Kessler, IDA, Linköpings universitet.

4

## Lectures (2)



...

- **Lecture 8:** Parallelization of sequential programs\*
- **Lecture 9:** GPU architecture and trends
- **Lecture 10:** Introduction to CUDA programming
- **Lecture 11:** CUDA programming. GPU lab introduction
- **Lecture 12:** Sorting on GPU. Advanced CUDA issues. FFT on GPU
- **Lecture 13:** Introduction to OpenCL.
- **Lesson 3:** OpenCL. Shader programming. Exercises.
- **Lecture 14:** High-level Parallel Programming with Skeletons\*
- **Lecture 15:** Advanced issues. Wrap-up

C. Kessler, IDA, Linköpings universitet.

5




## Outlook

Christoph Kessler, IDA,  
Linköpings universitet

## Related course

### TDDC78 Programming of Parallel Computers - Methods and Tools, 6hp, VT2

- Complement to TDDD56
- Focus on parallel scientific computing and clusters (NSC)
  - Parallel architecture concepts, including interconnection networks, shared and distributed memory architectures
  - Thread programming with OpenMP (Labs)
  - Message passing programming with MPI (Labs)
  - Performance analysis tools (Lab)
  - Parallel algorithms in scientific programming, including basic linear algebra and linear system solving
- Some minor overlap (→ synergy)
  - Shared memory architecture, Theory, Parallelizing loops, Skeleton programming



C. Kessler, IDA, Linköpings universitet

TDDD56 Programming parallel computers



## Master Thesis Topics Available!

Christoph Kessler, IDA,  
Linköpings universitet

## Master thesis projects available! (1)

### ■ Multicore programming techniques and tools for...

- GPU-based systems and GPU clusters
- Intel SCC 48-core research prototype
- Movidius Myriad1, Myriad2
- Planned: Adapteva Epiphany, ARM big.LITTLE, ...

■ ...






Image source: Movidius

C. Kessler, IDA, Linköpings universitet

Image source: www.kitray.eu

Image source: www.adapteva.com

## Master thesis projects available! (2)

### ■ Parallelization and optimization techniques

- Modeling and optimizing performance and energy efficiency of parallel software
  - ▶ Empirical modeling by adaptive sampling
- Mapping and scheduling tasks to parallel architectures
  - ▶ On-chip pipelining
  - ▶ Crown scheduling



### ■ Retargetability / Generic optimization

- Platform modeling
- Micro-benchmarking


C. Kessler, IDA, Linköpings universitet

10

## Example: On-Chip Pipelining

### ■ Mapping pipelineable streaming task graphs to manycore systems

- e.g. pipelined parallel mergesort



IBM Cell  
[J.UCS'12]

Intel SCC (48 cores)  
[ICCS-WEPA'12]

C. Kessler, IDA, Linköpings universitet

11

## Orchestrating cores and main memory for pipelineable (streaming) computations

### Standard model: Dancehall

- Use core-local memory as explicitly managed cache
- Bandwidth to main memory is performance bottleneck
- Easier to implement

### On-Chip Pipelining [K., Keller HPPC'08]

- All tasks active simultaneously
  - User-level scheduler
  - Round-robin, data driven
- Trades increased on-chip communication for reduced off-chip mem. access
- Example (Cell, PS3):  
Parallel merging 8M ints on 6 SPEs in 93 ms instead of 225 ms (CellSort)

Example:  
Merge tree computation





Image source: N. Melot

12


## Generalization

### Streaming computations

- Pipelining
- Concurrent execution of all streaming tasks in steady state

### Streaming task graphs

- (Acyclic) pipeline graph of streaming tasks
- Producer-consumer communication
- Cf. Kahn Process Networks



C. Kessler, IDA, Linköping universitet.

13


## Generalization

### Streaming computations

- Pipelining
- Concurrent execution of all streaming tasks in steady state


### Streaming task graphs

- (Acyclic) pipeline graph of streaming tasks
- Producer-consumer communication
- Cf. Kahn Process Networks



### Streaming task collections


- Independent streaming tasks
- OR: Steady state of the pipeline (one round): all tasks (instances) considered independent
- Computational load matters, ignore communication for now



C. Kessler, IDA, Linköping universitet.

14

## Mapping for the steady state




C. Kessler, IDA, Linköping universitet.

15

## Now: Moldable (Parallelizable) Tasks

### Moldable tasks $j = 1, \dots, n$

- A task  $j$  performs fixed **work**  $work_j$
- Can be run with (integer)  $w > 1$  cores
  - Resource allocation  $\rightarrow$  task **width**  $w_j$
  - Limit: **maximum width**  $W_j$ :  $w_j \leq W_j$  for task  $j$
- Arbitrary scalability functions: **efficiency**  $0 < e_j(w) \leq 1$  for  $1 \leq w \leq W_j$



### Mixed task model

Streaming tasks  $j = 1, \dots, n$  can be

- Inherently sequential (max. width  $W_j = 1$ )
- Malleable with limited scalability (with given  $W_j > 1$ )
- Malleable with unlimited scalability ( $W_j \geq p$ )


C. Kessler, IDA, Linköping universitet.

16

## "Streaming Task Scheduling"

### 3 static problems to solve (for the steady state of the pipeline):

- **Resource allocation:**
  - Allocate resources (width  $w_j \leq W_j$ ) to each task  $j$
  - Determines its parallel execution time
- **Mapping**
  - Map each task  $j$  to  $w_j$  specific cores for execution
- **Discrete Frequency Scaling**
  - Select for each task  $j$  a frequency  $f_j$  in  $\{f_1, \dots, f_s\}$  to match a throughput (round makespan) constraint
- Each core then executes its assigned (threads of) tasks at runtime round-robin in a data-driven manner
- Tasks that are not data ready are skipped for the current round



C. Kessler, IDA, Linköping universitet.


17

## Idea: Constrain Resource Allocation

- Define a hierarchy of processor groups

### The Crown

- Example: A balanced binary crown over 8 cores



- ▶ Core allocations (task widths)  $w_j$  must be powers of 2
- ▶ Reduces **#possible mapping targets** from  $2^p - 1$  to  $2p - 1$  (= #groups)

C. Kessler, IDA, Linköping universitet.

18

## Crown Scheduling

- Crown Allocation
- Crown Mapping
- Crown Scaling

**Crown Schedule:** one round of the steady state of the pipeline

- **Crown Schedule:** Same order and nonincreasing widths of tasks executed “down-crown” for each core
- Minimizes global interferences for parallel task scaling
- Opens for dynamic rescaling if tasks are dropped for a round

C. Kessler, IDA, Linköpings universitet. 19

## FFT Example, Solution by Stepwise CS

C. Kessler, IDA, Linköpings universitet. 20

## FFT Example, Solution by Integrated CS

C. Kessler, IDA, Linköpings universitet. 21

## Crown Scheduling

- Optimal and Heuristic Algorithms
  - Faster, and no worse schedule quality (energy) than competing approaches for moldable task scheduling
  - [Melot et al. ACM Trans. Arch. Code Opt., to appear 2015]
- Open problems ...
  - Consider communication load on on-chip network
  - Adapt for architectural properties e.g. voltage / frequency domains
  - Generic Crown Scheduler (read platform description, optimize, generate C code)
  - Automated microbenchmarking for moldable tasks
  - Graphical visualizer for schedules

C. Kessler, IDA, Linköpings universitet. 22

## Master thesis projects available! (3)

- Support for high-level parallel computation models
  - SkePU skeleton programming library
    - ▶ Energy tuning and adaptive sampling
    - ▶ Support for multi-variant user functions (SIMD, unrolling)
    - ▶ Streaming support for big-data computations
    - ▶ New target architectures, e.g. Myriad2
    - ▶ New skeletons, e.g. pipe, DC
    - ▶ New smart containers, e.g. image, sparse matrix
    - ▶ Cluster-SkePU
  - Multi-Variant Components with autotuned selection
    - ▶ Code variants
    - ▶ Data structures / layouts → smarter containers
    - ▶ Global composition strategies and framework

C. Kessler, IDA, Linköpings universitet. 23

## SkePU / Global Composition Framework

### Bulk Global vs. Greedy Local Selection

- Outperforming HEFT for chains of dependent component calls [Dastgeer, K., J. Supercomputing 2014, to appear]

ODE solver

C. Kessler, IDA, Linköpings universitet. 24

## Master thesis projects available! (4)



- **Application porting and algorithm engineering**  
case studies for multicore
  - Medical image processing / Stencil computations
  - Statistics kernels
  - /// IP forwarding data plane on GPUs  
(algorithms, portability, scheduling)
- Automatic extraction and transformation of parallelism
- ...

Join our research: [www.ida.liu.se/~chrke/exjobb](http://www.ida.liu.se/~chrke/exjobb)

C. Kessler, IDA, Linköpings universitet.

25

## The End (?)



*"Now, this is not the end.  
It is not even the beginning of the end.  
But it is, perhaps, the end of the beginning."  
-- W. Churchill*

C. Kessler, IDA, Linköpings universitet.

26

